

Disciplined locking: No more concurrency errors



<http://CheckerFramework.org/>

Werner Dietl, University of Waterloo
Michael Ernst, University of Washington



Concurrency: essential but error-prone

- + Essential for performance (exploit multiple cores)
- + Design component of GUIs
- Data races: concurrent access to shared data
 - easy mistake to make
 - leads to corrupted data structures
 - difficult to reproduce and diagnose

Thread-unsafe code

```
class BankAccount {  
  
    int balance;  
  
    void withdraw(int amount) {  
        int oldBal = this.balance;  
        int newBal = oldBal - amount;  
        this.balance = newBal;  
    }  
}
```

...

Data race example

Shared account
Initial balance = 500



Thread 1:

sharedAccount.withdraw(50)



Thread 2:

sharedAccount.withdraw(100)

```
int oldBal = this.500.ance;  
int newBal = o5001 - ar50nt;  
this.balance = n4501;
```

```
int oldBal = this.500.ance;  
int newBal = o5001 - a100nt;  
this.balance = n4001;
```

Withdrawals = 150
Final balance = 450

Solution: locking

```
class BankAccount {  
  
    Object acctLock;  
    @GuardedBy("acctLock") int balance;  
  
    void withdraw(int amount) {  
        synchronized (acctLock) {  
            int oldBal = this.balance;  
            int newBal = oldBal - amount;  
            this.balance = newBal;  
        }  
    }  
}
```

Locking:

- Only one thread can acquire the lock
- No concurrent access to data
- Which lock to hold?

Key issues:

- Names vs. values
- Aliasing

Locking discipline = which locks to hold when accessing what data

```
@GuardedBy("lock1") int w;  
@GuardedBy("lock2") int x;  
@GuardedBy("lock2") int y;  
                int z;
```

- Write locking discipline as documentation and for use by tools
- `@GuardedBy` [Goetz 2006] is a de-facto standard
 - On GitHub, 35,000 uses in 7,000 files
- Its semantics is **informal, ambiguous, and incorrect**
 - It allows data races
- Similar problems with other definitions

Outline

- Formal semantics for locking disciplines
 - unambiguous
 - **prevents data races** (other concurrency errors exist)
 - two variants: value-based, name-based
- Two implementations:
 - **type-checker that validates use of locking**
 - inference tool that infers locking discipline
- Experiments: programmer-written **@GuardedBy**:
 - are often inconsistent with informal semantics
 - permit data races even when consistent

Concurrency background

specification
of locking
discipline

synchronized
statement or
method locks
the monitor.

Exiting the
statement or
method unlocks
the monitor.

```
Date d = new Date();
```

```
@GuardedBy("d") List lst = ...;
```

```
synchronized (d){  
    lst.add(...)  
    lst.remove(...)  
    otherList = lst;  
}
```

Each object is
associated with a
monitor or *intrinsic lock*.

guard expression;
arbitrary, e.g. `a.b().f`

Our implementations
handle explicit locks too.

Defining a locking discipline

Guard expression:

- Aliases? Yes
- Reassignment? No
- Side effects? Yes
- Scoping? Def site

Value protection answers

Informally:

“If program element x is annotated by `@GuardedBy(L)`,
a thread may only use x
while holding the lock L.“

```
MyObject lock;  
@GuardedBy("lock.field") Pair shared;  
@GuardedBy("lock.field") Pair alias;
```

```
synchronized (lock.field) {  
    shared.a = 22;  
    alias = shared;  
}
```

Element being guarded:

- Name or value? Value
- Aliases? Yes
- Reassignments? Yes
- Side effects? Yes

What is a use?

- Occurrence of name?
- Dereference of name? (`x.f`)
- Dereference of value?
(Dereference = field read/write)

← current

← better

```
MyObject lock;  
@GuardedBy("lock") Pair shared;  
Pair alias;
```

Name protection

... not value protection

```
synchronized (lock) {  
    alias = shared;  
}  
  
alias.a = ...
```

Suffers a data race

Value protection

... not name protection

```
alias = shared;  
synchronized (lock) {  
    shared.a = ...  
}
```

No data race

alias.a = ...
is forbidden without
holding lock

Locking discipline semantics providing value protection

```
type qualifier  
@GuardedBy("lock") Pair shared;  
type variable
```

Suppose expression x has type: $\text{@GuardedBy}(L) C$

A *use* is a dereference

Type system constraint; may lock an alias

When the program dereferences a value that has ever been bound to x ,
the program holds the lock on the value of expression L .

The referent of L must not change while the thread holds the lock.

No reassignment of guard expression.

Side effects permitted (do not affect the monitor).

Locking discipline semantics providing name protection

```
@GuardedBy("lock") Pair shared;
```

The code is annotated with three blue curly braces. The first brace groups '@GuardedBy("lock")' as a 'variable annotation'. The second brace groups 'Pair' as the 'type'. The third brace groups 'shared' as the 'variable'.

Suppose variable v is declared as $\text{@GuardedBy}(L)$

A *use* is a variable read or write No aliasing permitted

When the program accesses v , which must not be aliased,
the program holds the lock on the value of expression L .
 L may only be “itself” or “this”.

Guarantees L always evaluates to the same value

Sound name protection
is restrictive

Demo

Key theoretical contributions

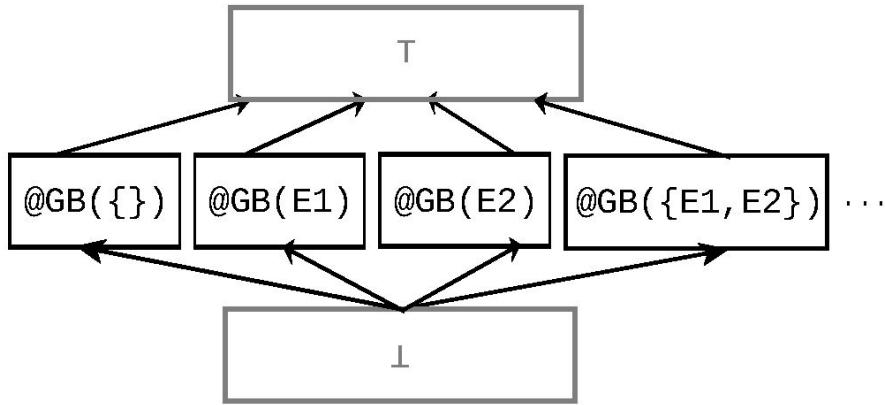
- Two formal semantics (name-based and value-based)
 - Core calculus based on RaceFreeJava [Abadi TOPLAS 2006]
 - Structural Operational Semantics
 - Definitions of accessed variables and dereferenced locations
- Proofs of correctness
 - By contradiction:
 - assume data race
 - show locking discipline must have been violated

Static analysis of a locking discipline

- Goal is to determine facts about **values**
 - Program is written in terms of facts about **variables**
- Analysis computes an approximation (an abstraction)
 - of values each expression may evaluate to
 - of locks currently held by the program

Both abstractions
are sound

Enforcement of value semantics via type-checking



Type rule:
If $x : @\text{GB}(L)$,
then L must be held
when x is dereferenced

- No two `@GuardedBy` annotations are related by subtyping
- Why not $@\text{GB}(L1) <: @\text{GB}(L1, L2)$?
 - Side effects and aliasing

More type system features

- Method pre/postconditions (@Holding annotations)
- Side effect annotations
- Type qualifier polymorphism
- Reflection
- Flow-sensitive type inference

Type-checking is a modular analysis

Modular = one procedure at a time

When type-checking a method, examine *only*:

- its **body**
- the **signatures of its callees**

Advantages: fast; clear error messages

Disadvantage: need library annotations

Is this code safe?

```
@GuardedBy("myMonitor") List<String> myList;  
System.identityHashCode(myList);
```

Is this code safe?

```
@GuardedBy("myMonitor") List<String> myList;  
System.identityHashCode(myList);  
  
myList.clear();
```

clear() requires synchronization

```
@GuardedBy("myMonitor") List<String> myList;  
System.identityHashCode(myList);  
synchronized(myMonitor) {  
    myList.clear();  
}
```

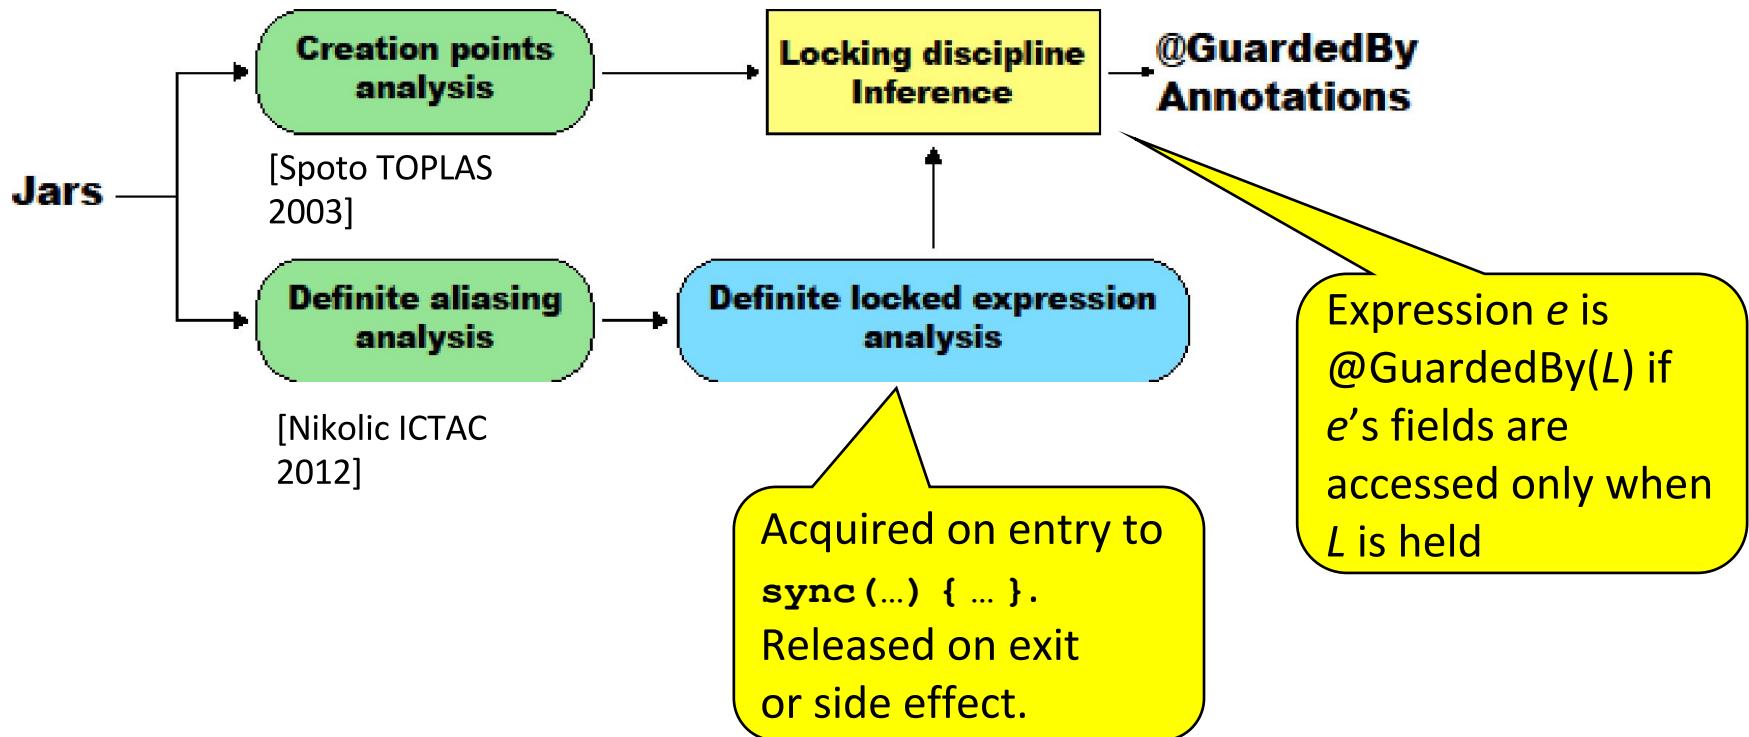
Library annotations = constraints on callers

```
@GuardedBy("myMonitor") List<String> myList;  
System.identityHashCode(myList);  
synchronized(myMonitor) {  
    myList.clear();  
}
```

```
class System {  
    @ReleasesNoLocks  
    static int  
    identityHashCode(Object x);  
}
```

```
interface List<E> {  
    @ReleasesNoLocks  
    boolean  
    clear(@GuardSatisfied List<E> this);  
}
```

Inference of both semantics via abstract interpretation



Inference implementation

1. Where is the guarded element used?
 - Name protection: syntactic uses of variable
 - Value protection: estimate via creation points analysis
2. What expressions are locked at those points?
 - Definite aliasing analysis
 - Side effect analysis
 - Viewpoint adaptation (contextualization)

Whole-program analysis

- Makes closed-world assumption
- Type-checking is modular, incremental

Experimental evaluation of value semantics

- 15 programs, 1.3 MLOC
 - BitcoinJ, Daikon, Derby, Eclipse, Guava, Jetty, Velocity, Zookeeper, Tomcat, ...
 - 5 contain programmer-written `@GuardedBy` annotations
- 661 correct annotations
 - Candidates: annotations written by the programmer or inferred by our tool
 - Correct: program never suffers a data race on the element
 - Determined by manual analysis

Experimental results

Are the annotations
correct?

Are the annotations
complete?

| | Precision | Recall |
|---------------|-----------|--------|
| Inference | 100% | 83% |
| Type-checking | 100% | 99% |
| Programmers | 50% | 42% |



Tell the whole truth and nothing but the truth.

Programmer mistakes

Errors in every program that programmers annotated with respect to both value and name semantics

- Wrong lock
 - Guava: `@GuardedBy("Segment.this")` should be `@GuardedBy("this")`
- Lock writes but not reads
 - Guava: `SerializingExecutor.isThreadScheduled`
- Omitted annotations
 - BitcoinJ: `PaymentChannelServer` ; Apache Velocity: several
- **Creating external aliases**
 - BitcoinJ: `PaymentChannelClient.conn`

Implementations

- Type checker:

- Lock Checker, distributed with the Checker Framework
- <http://CheckerFramework.org/>
- Live demo: <http://eisop.uwaterloo.ca/live>



- Inference:

- Julia abstract interpretation
- <http://juliasoft.com/>



The Checker Framework

A framework for pluggable type checkers

“Plugs” into the OpenJDK or OracleJDK compiler

```
javac -processor MyChecker ...
```

Standard error format allows tool integration



Eclipse plug-in

```
3 public class Test {  
4  
5     public static void main(String[] args) {  
6         Console c = System.console();  
7         c.printf("Test");  
8     }  
9 }
```

Problems @ Javadoc Declaration Search

0 errors, 1 warning, 0 others

Description

Warnings (1 item)

dereference of possibly-null reference c
c.printf("Test");

```
3 public class Test {  
4  
5     public static void main(String[] args) {  
6         Console c = System.console();  
7         dereference of possibly-null reference c.c.printf("Test");  
8     }  
9 }
```

Problems @ Javadoc Declaration Search Console Task

0 errors, 1 warning, 0 others

Description

Warnings (1 item)

dereference of possibly-null reference c
c.printf("Test");

Resource

Test.java



Ant and Maven integration

```
<presetdef name="jsr308.javac">
  <javac fork="yes"
    executable="${checkerframework}/checker/bin/${cfJavac}" >
    <!-- JSR-308-related compiler arguments -->
    <compilerarg value="-version"/>
    <compilerarg value="-implicit:class"/>
  </javac>
</presetdef>
```

```
<dependencies>
  ... existing <dependency> items ...
  <!-- annotations from the Checker Framework:
      nullness, interning, locking, ... -->
  <dependency>
    <groupId>org.checkerframework</groupId>
    <artifactId>checker-qual</artifactId>
    <version>1.9.7</version>
  </dependency>
</dependencies>
```

Live demo <http://eisop.uwaterloo.ca/live/>

Checker Framework Live Demo

Write Java code here:

```
1 import org.checkerframework.checker.nullness.qual.Nullable;
2 class YourClassNameHere {
3     void foo(Object nn, @Nullable Object nbl) {
4         nn.toString(); // OK
5         nbl.toString(); // Error
6     }
7 }
```

Choose a type system: ▾

Examples:

Nullness: [NullnessExample](#) | [NullnessExampleWithWarnings](#)

MapKey: [MapKeyExampleWithWarnings](#)

Interning: [InterningExample](#) | [InterningExampleWithWarnings](#)

Lock: [GuardedByExampleWithWarnings](#) | [HoldingExampleWithWarnings](#) | [EnsuresLockHeldExample](#) | [Loc](#)



Example type systems

Null dereferences (@NotNull)

>200 errors in Google Collections, javac, ...

Equality tests (@Interned)

>200 problems in Xerces, Lucene, ...

Concurrency / locking (@GuardedBy)

>500 errors in BitcoinJ, Derby, Guava, Tomcat, ...

Fake enumerations / typedefs (@Fenum)

problems in Swing, JabRef



Checkers are usable

- Type-checking is **familiar** to programmers
- Modular: fast, incremental, partial programs
- Annotations are **not too verbose**
 - **@NonNull**: 1 per 75 lines
 - **@Interned**: 124 annotations in 220 KLOC revealed 11 bugs
 - **@Format**: 107 annotations in 2.8 MLOC revealed 104 bugs
 - Possible to annotate part of program
 - Fewer annotations in new code
- Few false positives
- First-year CS majors preferred using checkers to not
- **Practical**: in daily use at Google, on Wall Street, etc.



Related work

- Name-based semantics: JML, JCIP, rccjava [Abadi TOPLAS 2006], ...
- Heuristic checking tools: Warlock, ESC/Modula-3, ESC/Java
- Unsound inference: [Naik PLDI 2006] uses may-alias, [Rose CSJP 2004] is dynamic
- Sound inference for part of Java [Flanagan SAS 2004]
- Type-and-effect type systems: heavier-weight, detects deadlocks too
- Ownership types

Learning more

Using Type Annotations to Improve Your Code

BoF3427, tonight, 19:00 to 19:45

Continental Ballroom 4

Technical papers:

- “Locking discipline inference and checking”, ICSE 2016
- “Semantics for locking specifications”, NFM 2016

Try the Lock Checker: <http://CheckerFramework.org/>



How to prevent (some) concurrency errors

- Formal semantics for locking disciplines
 - unambiguous, prevents data races
- Experiments: programmer-written `@GuardedBy`:
 - are often inconsistent with informal semantics
 - permit data races even when consistent with informal semantics
- Type-checker validates use of locking discipline
`(@GuardedBy)`
 - Download from <http://CheckerFramework.org/>